

CSCI 3753

Operating Systems

Introduction

Lecture Notes By
Shivakant Mishra
Computer Science, CU-Boulder
Last Update: 01/09/13

Operating Systems

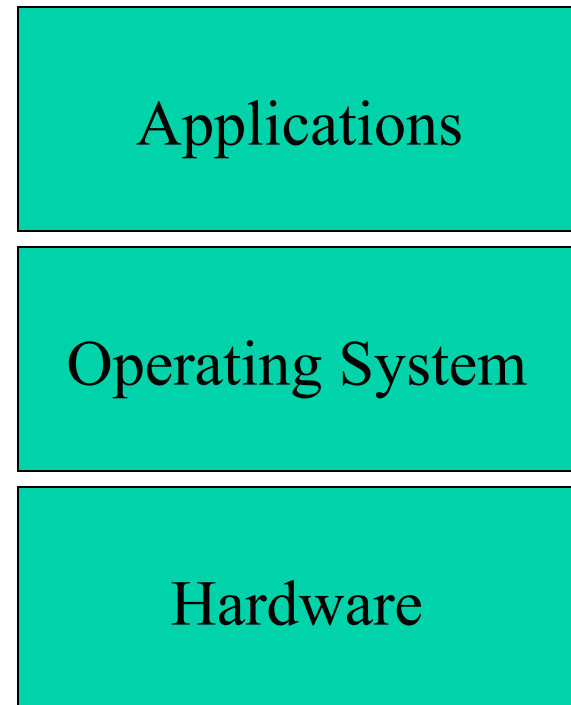
- OS is an essential (and perhaps the most important) part of a computing system.
 - PCs, laptops, supercomputers, cloud, data centers, cell phones, PDAs, tablets, embedded devices, ...
 - Controls hardware components.
 - Provides a usable interface.
 - Allows sharing of various computing components.
- One of the largest piece of software.

Windows, Linux, Mac OS X, Google Android, ...

> 500 at http://www.operating-system.org/betriebssystem/_english/os-liste.htm

What is an Operating System?

An operating system is a layer of software between applications and hardware that provides useful services to applications




Two views of an OS

- Extended machine
 - Hardware is too complex for most computer users to understand.
 - OS provides users with an equivalent of an extended or virtual machine that is easier to use.
- Resource manager
 - A computing system is comprised of many resources: processors (CPUs), memory, clocks, disks, key board, mouse, monitor, network cards, etc.
 - OS allows sharing and effective utilization of these resources.

Security and protection

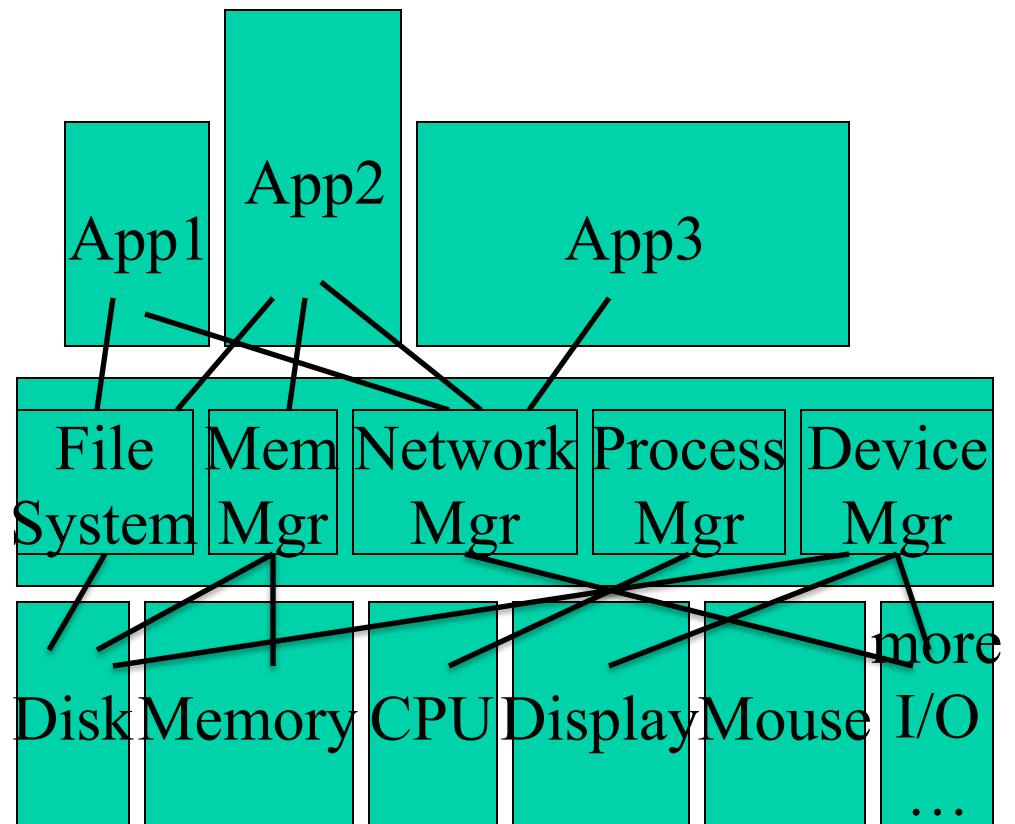
What is an Operating System?



Other devices include: wired network card, WiFi, camera, microphone, audio output, keyboard, DVD/CD, USB, etc.

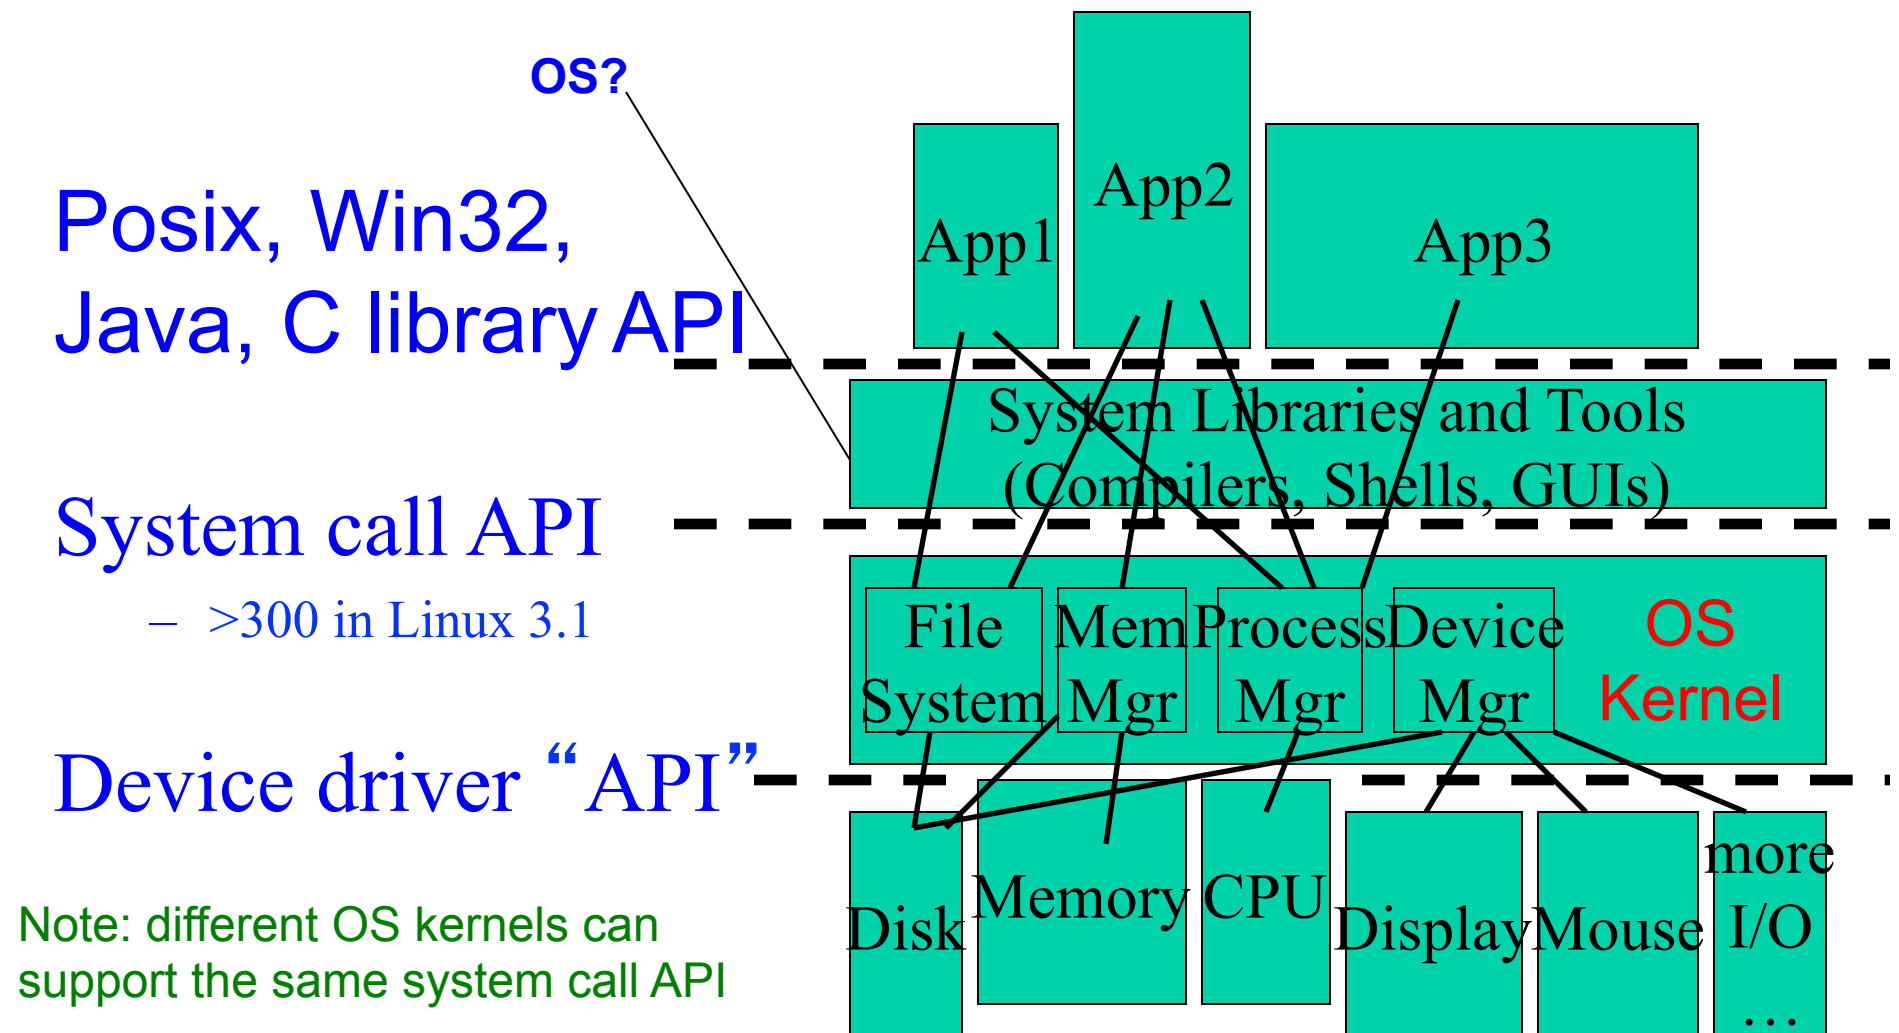
OS as Resource Manager

- Process management
- Memory management
- File system
- Device management
- Network management



Not pictured above in the OS:
the network manager.

Extended Machine View




Evolution of OS

- 1940' s and 50' s
 - Vacuum tubes
 - Single user
 - No programming language: machine language
 - Used for straight-forward numerical calculations
 - Single program
 - No OS
- mid 50' s and early 60' s
 - Transistors (more reliable, smaller in size, cheaper)
 - Punch cards
 - **Batch processing**
 - Used for scientific and engineering calculations.
 - FORTRAN

Batch Processing

- Execute a pre-defined collection of programs (jobs) called a batch.
- No human interaction

Batch Processing

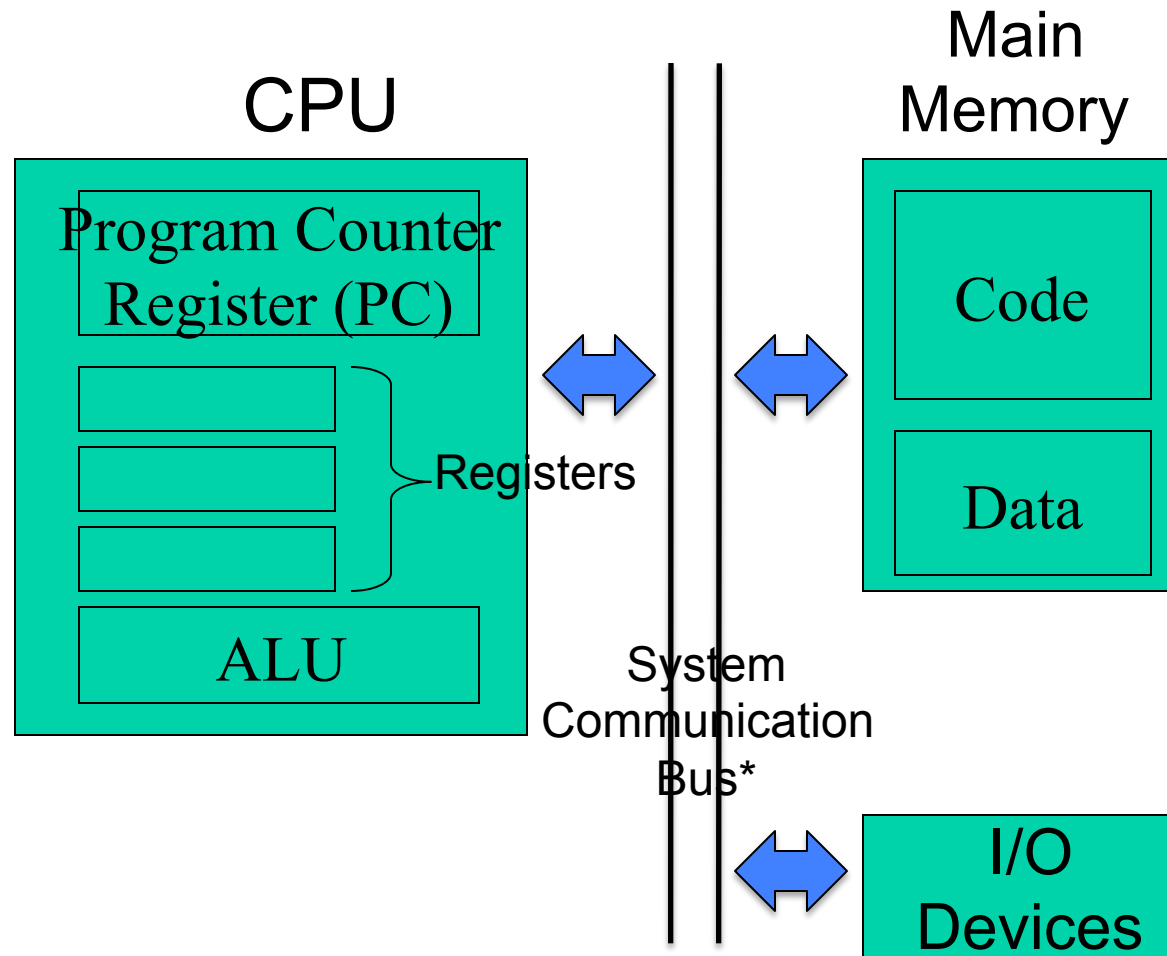


Limitations of Batch Processing

- Batch jobs are very non-interactive
 - Don't support a shell application for example
 - design jobs to yield much sooner than an I/O block, to give the impression of interactivity
 - If one of the programs was a shell, then it appears to the human user as if the computer is instantly responsive
 - In the small time segment a shell is given, it can draw a character on the screen that you've just typed => appearance of real-time interactivity


- mid 60' s to mid 70' s
 - Integrated circuits.
 - **Multiprogramming:** When CPU is idle, e.g. when the running program is blocked for an I/O, start another program.
 - **Multitasking:** Switch CPU between different programs irrespective of whether the running program is blocked.
 - Examples: CTSS, IBM 360, Multics, Unix

The von Neumann Computer's Hardware Architecture and I/O

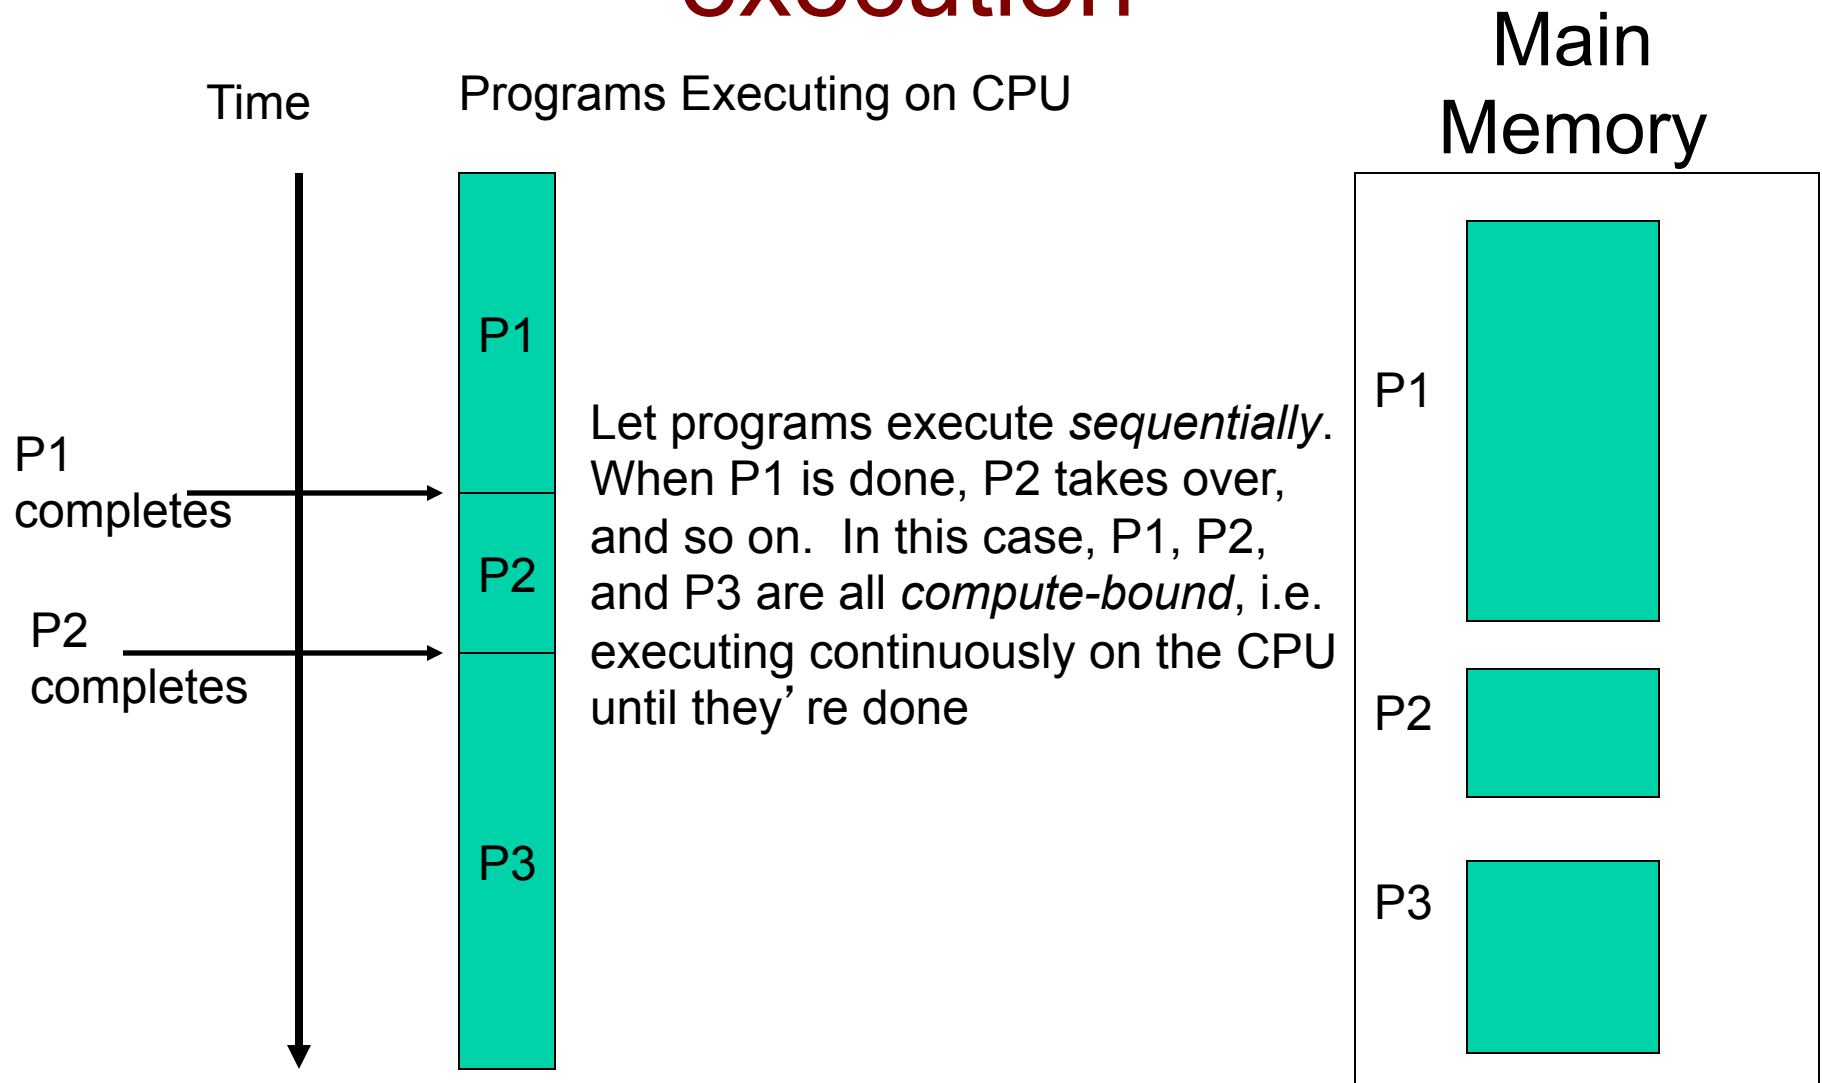


* Includes control, address, and data buses

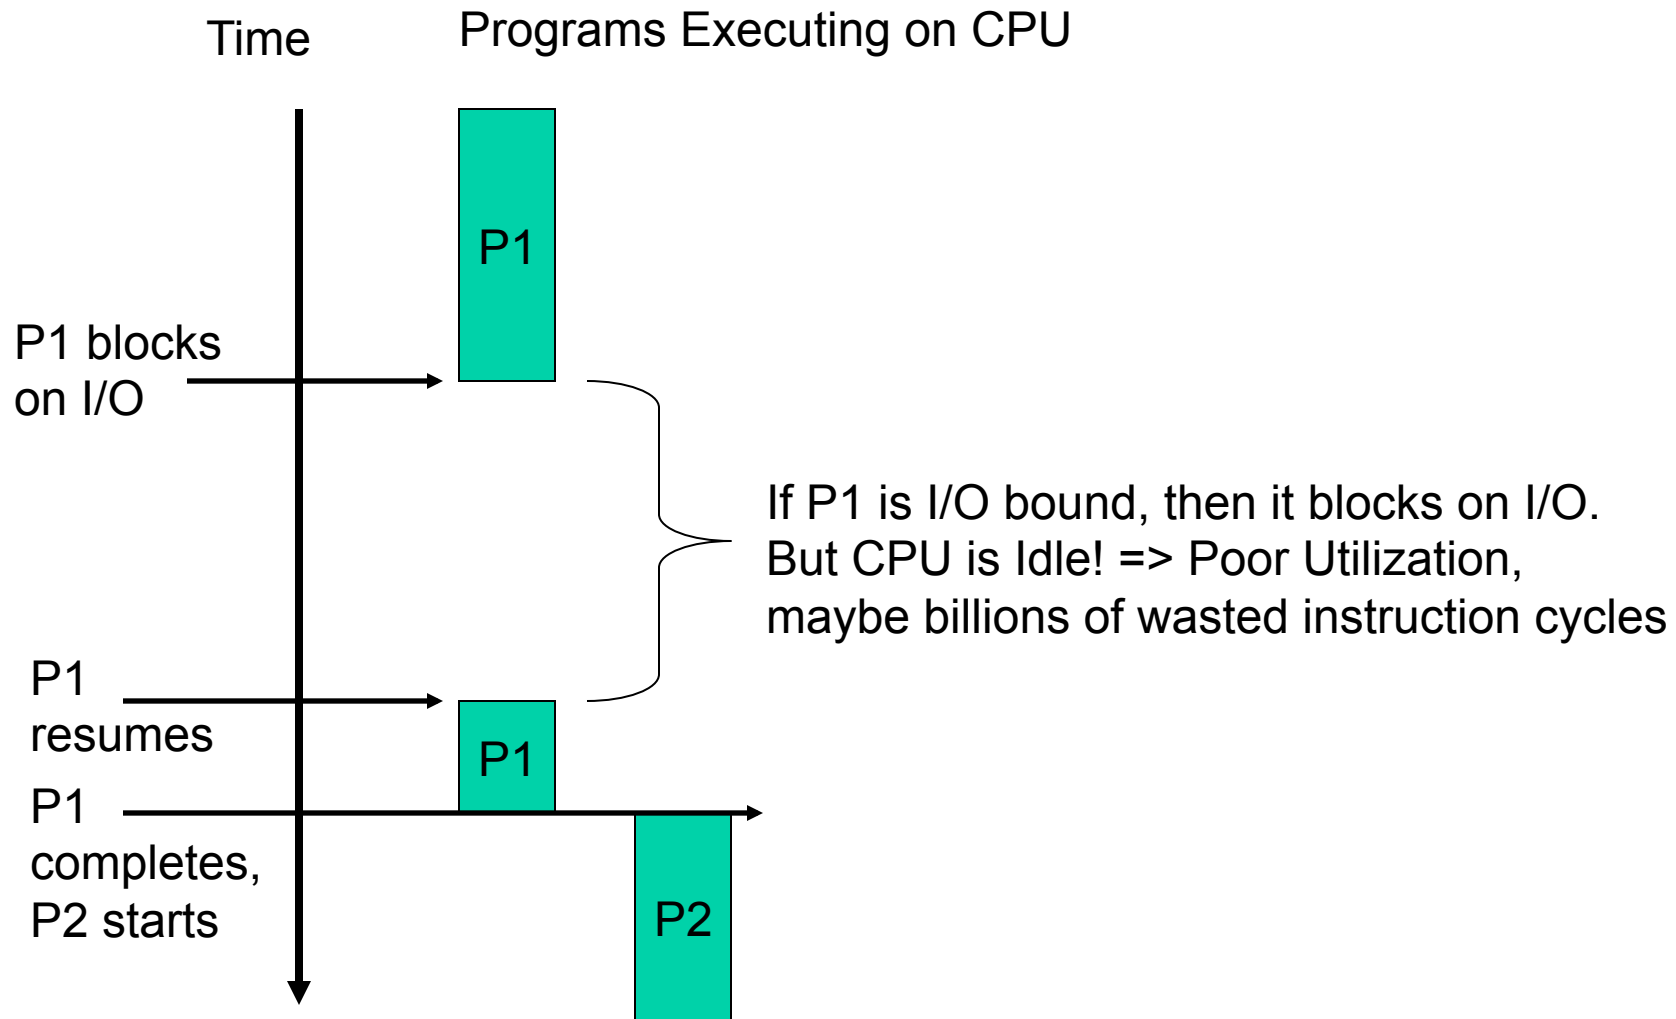
A Typical PC Bus Structure



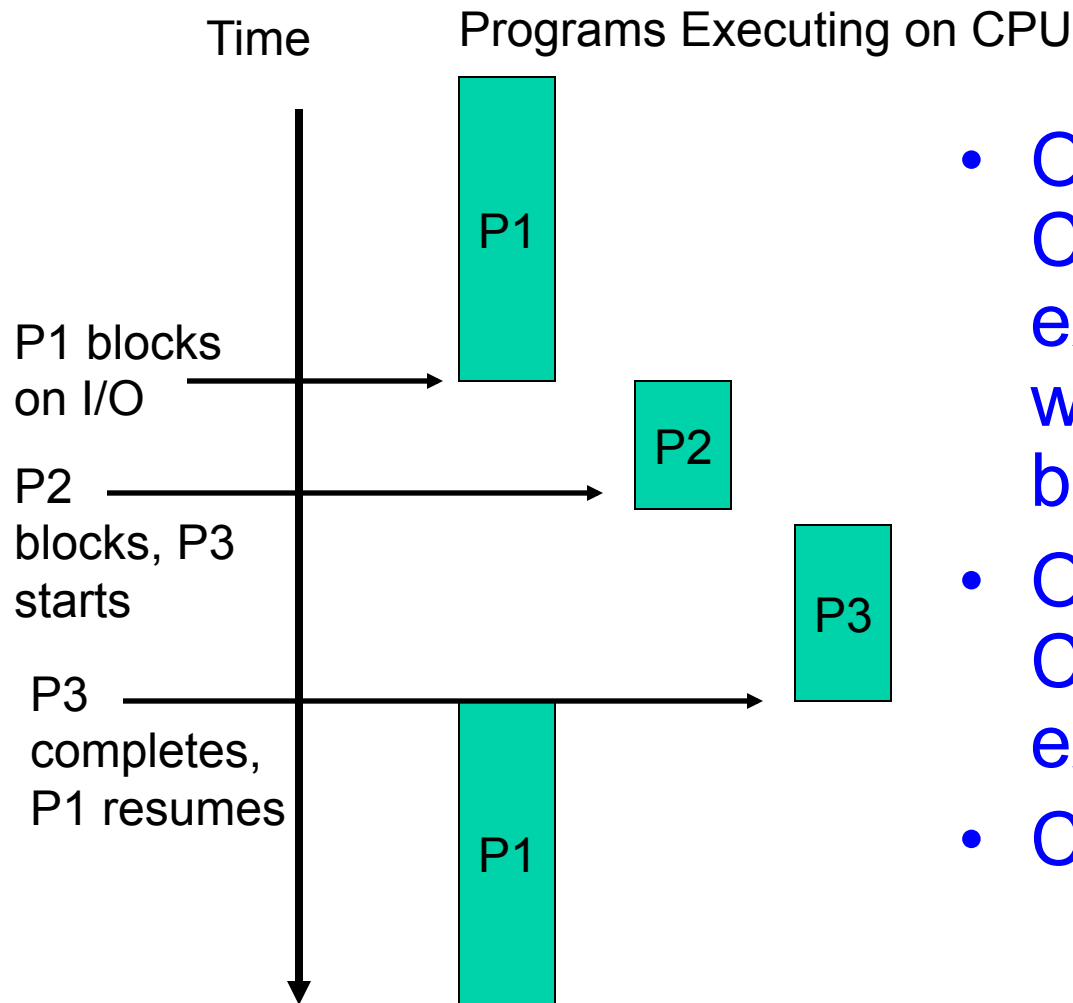
Multiprogramming: sequential execution



Multiprogramming: Problem with sequential execution



Multiprogramming



- OS Scheduler switches CPU between multiple executing programs when a program is blocked, e.g. for I/O
- OS *time-multiplexes* CPU between executable programs
- Context switch

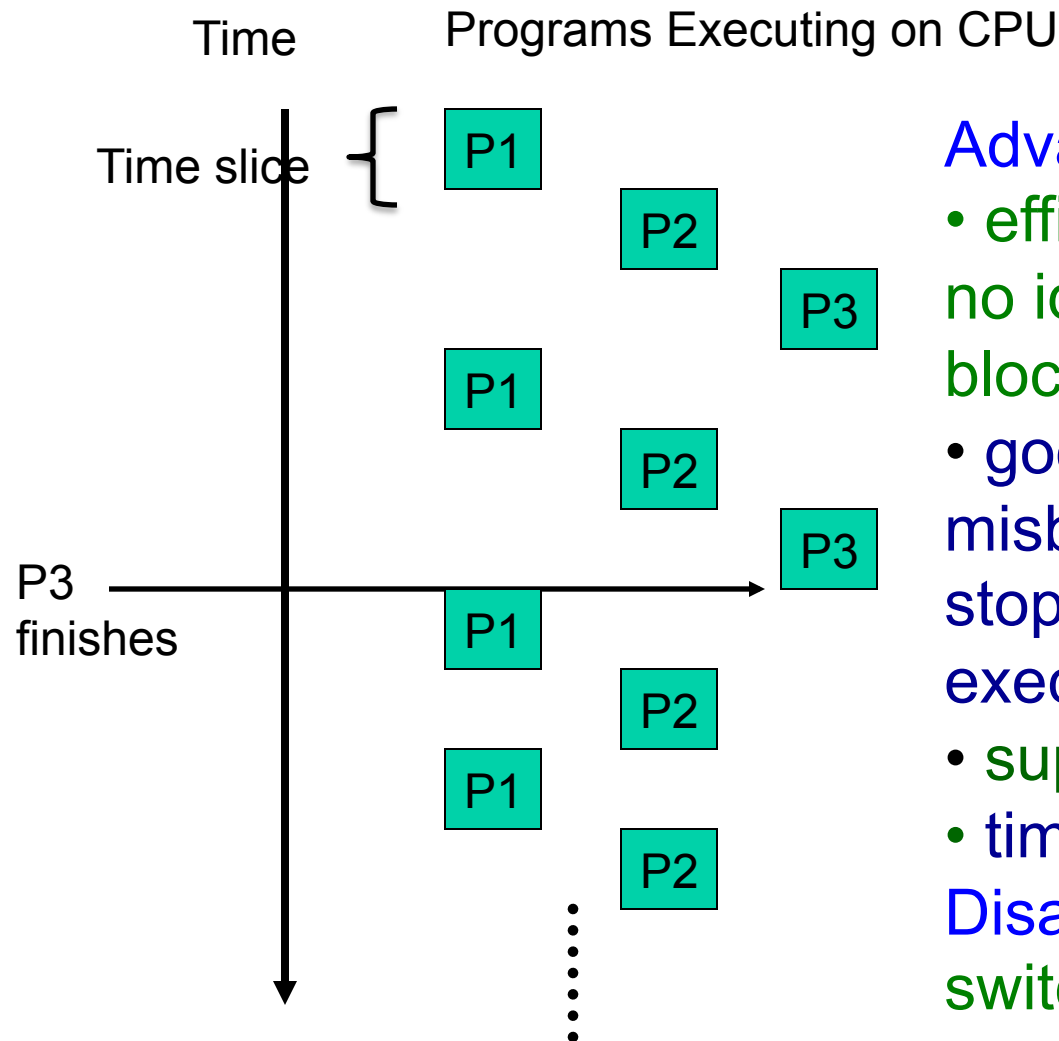
Context Switch

- Switching CPU from one running program to another is called a *context switch*
 - Save the current state (PC, IR, data registers, stack pointers, etc.) of the running application
 - Load the state of the new application
 - there is overhead due to this context switching
- No useful progress occurs for any applications while the OS is context-switching the CPU.

- Problem with pure multiprocessing: A CPU-bound program may delay the execution of other programs
 - A program with an infinite loop
 - A program to calculate the value of pi to the one-billionth decimal place

Multitasking

- CPU rapidly switches between programs



Advantages:

- efficient CPU usage, i.e. no idle time if one program blocks
- good isolation – a misbehaving program can't stop other programs from executing
- supports interactivity
- timesharing

Disadvantage: context switch overhead

Cooperative vs Preemptive Multitasking

- In cooperative multitasking, programs quickly and voluntarily yield CPU before they're done
 - Early OSs did this (Windows 3.1, Mac OS 9.*)
 - Poor fault isolation due to misbehaving programs
- In preemptive multitasking, OS forces programs to give up CPU
 - Fault isolation, interactivity, efficient CPU utilization
 - Time slice: time interval for which a program is allowed to run at any time
- All modern OSs are preemptively multitasked
 - Linux, BSD Unix, Windows NT/XP/Vista/7, Mac OS X 10.*

Preemptive Multitasking

- How does the OS force *rapid* switching?
 - Timer interrupt fires periodically
 - This suspends execution of the currently executing program and returns control to the OS scheduler
 - The scheduler decides the next program to execute and loads it, then passes control to it
- Context switch overhead: If you choose your default time slice too small, then you'll incur a lot of context-switching overhead as a % of your overall CPU utilization.


- mid 70' s to 90' s
 - LSI/VLSI; LANs
 - PCs and workstations
 - Process control and real-time systems
 - User friendly software
 - Distributed operating systems
 - MS DOS, 4.2 BSD Unix with TCP/IP, Mac OS with GUI, Linux, ...

OS Design

- OS design is guided by two factors:
 - Technology: CPU, storage, communication, new I/O devices, ...
 - Application needs

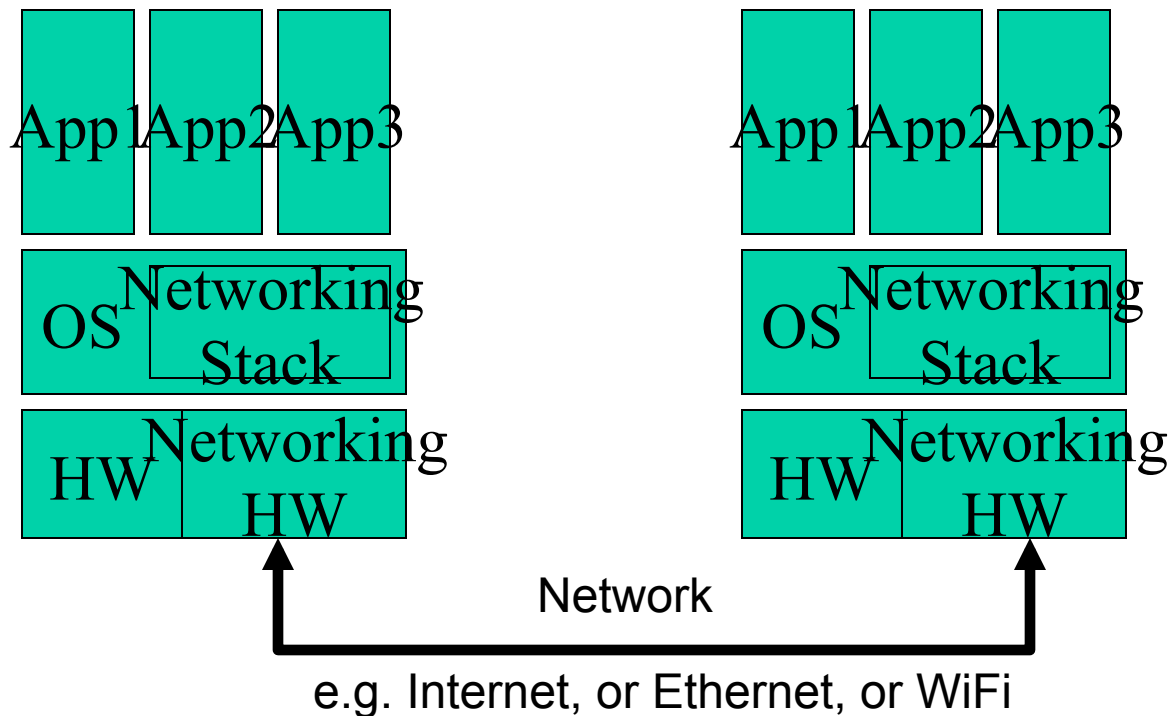
Operating System: Current Trends

- Hardware has evolved quickly - OS must adjust
 - Moore's Law roughly applies to CPU speed and/or memory size: doubles every 18 months => exponential!
 - Enables complex modern operating systems like Linux, Windows, UNIX, OS X



But Moore's Law doesn't apply to disk access speed or to battery life

Distributed operating systems



- Networked File System
- OS adds TCP/IP Network Stack
- Device driver support for Networking cards

- Examples:
 - App1 is a distributed client server app, e.g. App1 on left is Web browser, App1 on right is Web server

Operating System: Current Trends

- Diversification of OS' s to many different target environments
 - Energy-efficient cell phone OSs - scaling down
 - iPhone' s iOS, Google' s Android, ...
 - Multi-processor OSs - scaling up
 - Adapting Linux and Windows to multiple cores. Massively parallel supercomputers.
 - Real-Time OS for Embedded and Multimedia Systems
 - VXWorks, robotic OSs, ...

Operating System: Current Trends

- Virtualization – Virtual Machines (VMs)
 - Running a Windows VM inside a Linux OS, and vice versa.
 - More layers of abstraction
- Cloud computing rents VMs on racks of PCs at a massive scale

Google Data Center in The Dalles, Oregon

Size of
football
field

